

Analysis of Massive Data Sets

<http://www.fer.hr/predmet/avsp>

Prof.dr.sc. Siniša Srbljić

Doc.dr.sc. Dejan Škvorc

Doc.dr.sc. Ante Đerek

Faculty of Electrical Engineering and Computing
Consumer Computing Laboratory

Analysis of Massive Data Sets

Detection of near-duplicate documents using locality sensitive hashing

Otkrivanje sličnih dokumenata koristeći sažimanje neosjetljivo na lokalne promjene

Klemo Vladimir

Outline

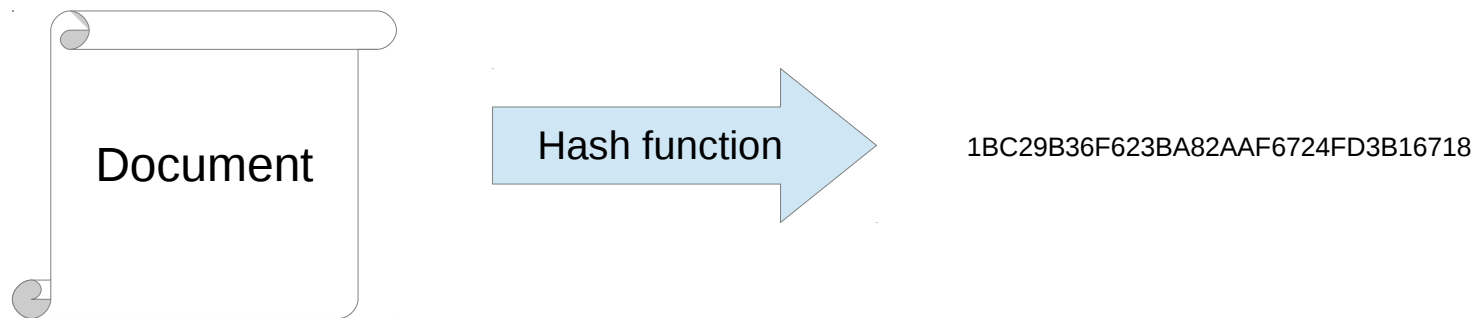
- Motivation
- Fingerprinting
 - Simhash algorithm
 - Rabin's rolling hash function
- Scalable queries
 - MapReduce sketch
- Literature

Motivation

- Near-duplicate documents
 - Versioning
 - Different versions of single document
 - Revisions, different file formats, ...
 - Mirroring
 - published in more than one location
 - Plagiarism
 - Exact or “processed” copy
 - Malware
 - Viruses, spam, ...
- Scalability
 - Documents and document repositories are large

Exact copy analysis

- Checksumming
 - Cryptographic hash functions
 - MD5, SHA1, SHA2, ...



- Catches the smallest edit
 - Great for detection of **exact** copies
 - Not so good for near-duplicate detection
 - Even the smallest change will result in totally different digest

Near-duplicate detection

- Two families of methods
 - 1. *Fingerprinting*
 - Document hashing ^{*1}
 - Dimensionality reduction
 - 2. Ranking
 - Information retrieval techniques ^{*2}
 - High-dimensional vectors manipulation

Fingerprinting

- Similarity preserving hashing
 - X , set of inputs
 - d_x , distance function on X
 - x_1, x_2 elements of X
 - similarity preserving hash function
 - $h: X \rightarrow Y$
 - $|Y| < |X|$
 - d_y , distance function on Y

Fingerprinting

- Similarity preserving hashing
 - Similar inputs have similar hashes

if $d_x(x_1, x_2) < \varepsilon_x$, then

$$d_y(h(x_1), h(x_2)) < \varepsilon_y$$

- Illustrative example

– $h("1234") = 0xaaaf$ $h("5678") = 0xb115$


– $h("1234") = 0xaaaf$ $h("1235") = 0xaaae$

Fingerprinting

- **Simhash** algorithm ^{*3}
 - Author: M. Charikar, 2002
 - Fingerprinting technique
 - Fingerprints of near-duplicates differ in a small number of bit positions (**hamming** distance)
 - Dimensionality reduction
 - Maps high-dimensional vectors to small-sized fingerprints (f-bits)
 - Input
 - High-dimensional vector (strings, numbers, ...)
 - Output
 - f-bit fingerprint
 - Fingerprint size (f-bits)
 - f is small and arbitrary
 - eg. f=64 for web sites

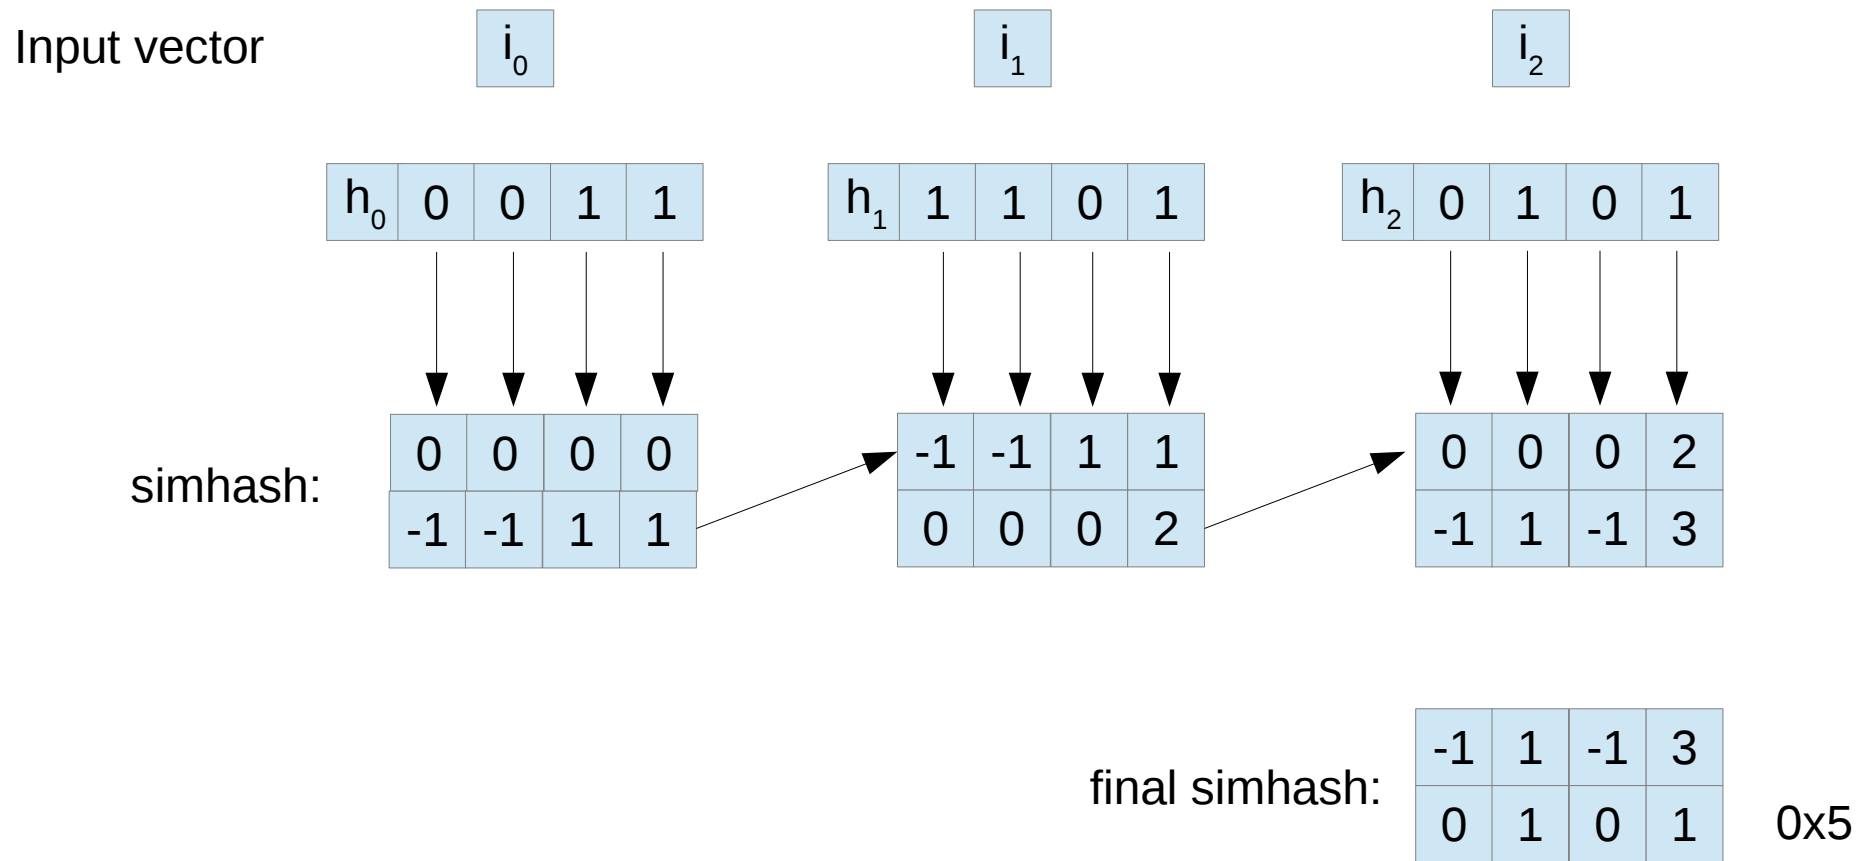
Simhash

- Simhash computation, $f=4$




Simhash

- Simhash example, $f=4$



Simhash

- Weighted Simhash computation
 - Assign weight factor to each feature



if $h_{yx} == 1$: $sh_x += w_y$

if $h_{yx} == 0$: $sh_x -= w_y$

Simhash

- Choice of hash function h
 - Uniform distribution
 - Fast
 - Candidates
 - Cryptographic hash functions
 - MD5 (128-bit), SHA-1 (256-bit)
 - Problem: cryptographic hashing is slow
 - Rolling hash functions
 - Input is hashed by moving window element by element

Simhash

- Choice of hash function h
 - Rabin rolling hash
 - Used in Rabin-Karp string searching algorithm
 - Input: string
 - $h(k) = k \bmod q$, q is some large prime number
 - Computation
 - substring coded as a number with base d
 - d = total number of possible characters
 - Coded string at position i :

$$x_i = s[i]d^{k-1} + s[i+1]d^{k-2} + \dots + s[i+k-1]$$

Simhash

- Choice of hash function h
 - Rabin rolling hash

- Example: $k=4$, $d=32$

- $i=0$

0	1	2	3	4		5	6	7	8	9
l	o	r	e	m		i	p	s	u	m
l	o	r	e							

$$h(\text{"lore"}) = h(x_0) = x_0 \bmod q$$

$$x_0 = \text{int}('l')32^3 + \text{int}('o')32^2 + \text{int}('r')32 + \text{int}('e')$$

$$x_{i+1} = ?$$

Simhash

- Choice of hash function h
 - Rabin rolling hash
 - Example: $k=4$, $d=32$
 - $i=1$

0	1	2	3	4		5	6	7	8	9
l	o	r	e	m		i	p	s	u	m
	o	r	e	m						

$$h(\text{"orem"}) = h(x_1) = x_1 \bmod q$$

$$x_1 = (x_0 - \text{int}('l')d^3)d + \text{int}('m')$$

$$\mathbf{x_{i+1} = (x_i - s[i]d^{k-1})d + s[i+k]}$$

Simhash

- Choice of hash function h
 - Rabin rolling hash
 - Fingerprints represented using polynomials
 - Computes hash value of the next string from the previous one
 - Constant number of operations
 - Independent of string length

Simhash

- Input vector
 - Focus on raw text documents
 - Convert document to a feature vector
 - Feature extraction
 - Tokenization
 - Unigram, 2-gram, 3-gram, ...
 - Stemming
 - Stopword removal
 - Phrase detection

Simhash

- Tokenization examples

`"lorem ipsum dolor sit amet"`

1. word tokens

`"lorem", "ipsum", "dolor", "sit", "amet"`

2. 2-word tokens

`"lorem ipsum", "ipsum dolor", "dolor sit", "sit amet"`

3. character 3-grams

`"lor", "ore", "rem", "em ", "m i", ...`

Simhash

- Shingle
 - hash of k-gram
 - k-grams
 - Characters, words, phrases, sentences (or combination)
 - $k = ?$
 - Small k : dissimilar documents appear similar
 - Large k : similar documents appear dissimilar
 - Feature vector from IR output
 - Weighted (TF) with IDF (inverse document frequency)
 - Might change when collection changes!

Fast Queries


- F , collection of f -bit fingerprints
- Q , query
 - single or set of fingerprints
- Task
 - identify whether Q differs from any of the fingerprints in F in at most k bits

Fast Queries

- Google numbers
 - 8B 64-bit fingerprints = 64GB
 - Online query
 - Q = single fingerprint
 - Restriction: few milliseconds
 - Batch query
 - Q = set of fingerprints
 - e.g. $|Q| = 1\text{M}$
 - Restriction: $\sim 100\text{seconds}$
 - 1B queries per day

Simhash

- Distribution of fingerprints
 - 8B 64-bit fingerprints



Fast Queries

- 1. approach
 - Build sorted table of **F**
 - Build list **Q'** with all fingerprints whose Hamming distance from **Q** is at most **k**

Q
100

k=1

Q'		
1	0	0
1	0	1
1	1	0
0	0	0

F		
0	0	0
0	0	1
0	1	0
0	1	1
1	0	0
1	0	1
1	1	0
1	1	1

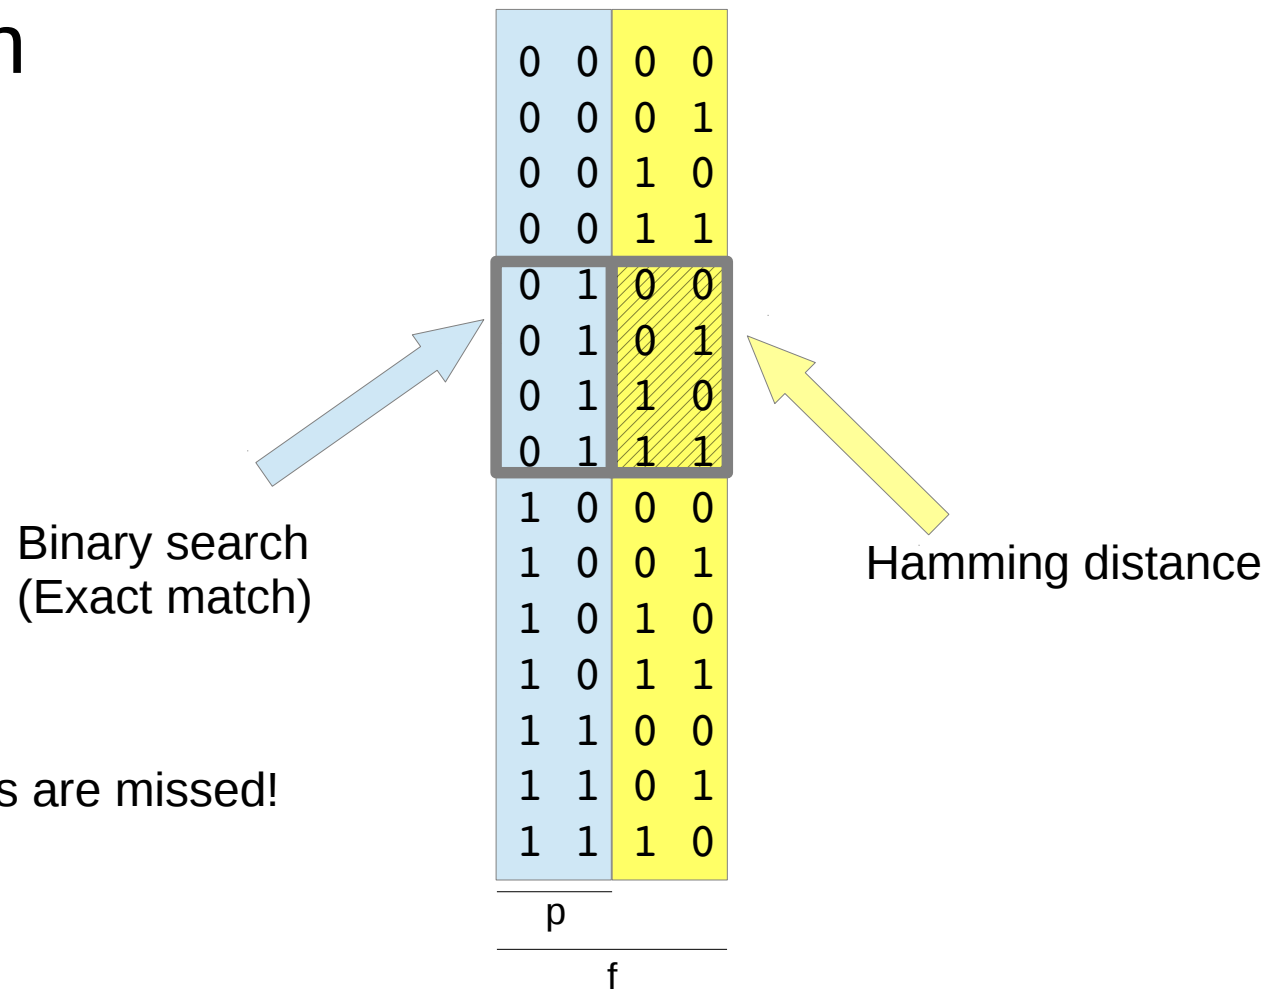
- 8B 64-bit fingerprints (k=3)
 - $|Q'| = \binom{64}{3} = 41664$

Fast Queries

- 2. approach
 - 1. Build sorted table of \mathbf{F}
 - 2. Find set of fingerprints (\mathbf{F}') that have equal most significant part (p bits)
 - Sorted table – binary search $O(p)$
 - 3. Check Hamming distance for each fingerprint in \mathbf{F}'
 - This approach will locate all fingerprints in \mathbf{F} that differ in at most k bits
 - **Restricted to least significant f-p bits!**

Fast Queries


- Illustration



Since $p < f$:
some true positives are missed!

Fast Queries

- Increasing precision/recall
 - Generate additional table
 - Reversed positions





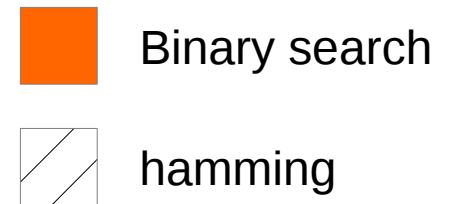



Fast Queries

- Solution
 - Build additional tables
 - Each with *different permutation* of bits
 - Every table has different set of significant bits
- Algorithm for fast (online) queries
 - Build t sorted tables of fingerprints: T_1, T_2, \dots, T_t
 - Each table T_i also contains
 - p_i – number of significant bits
 - Π_i – random permutation
 - *Every fingerprint in T_i is permuted with permutation Π_i*

Fast Queries

- Illustration
 - Build tables
 - Query in parallel



Fast Queries

- For given Q and k
 - Read each table (in parallel)
 - 1. Get fingerprints in T_i whose significant p_i bits match the significant p_i bits of $\Pi_i(Q)$
 - T'_i
 - $O(p_i)$ steps (binary search)
 - 2. For each fingerprint in T'_i , check if it's Hamming distance is at most k bits from $\Pi_i(Q)$

Fast Queries

- Example with $t=20$, $f=64$, $k=3$, $|F| = 8B$ (2^{34})
 - Split f into 6 blocks ($4 \times 11 + 2 \times 10$ bits)
 - Select 3 out of 6 blocks $\binom{6}{3} = 20$ ways
 - Arrange those blocks as significant bits
 - p = sum of those bits
 - 31, 32, or 33
 - On average query returns $2^{34-31}=8$ fingerprints

Fast Queries

- t and p parameters
 - $t \sim p$
 - Query time $\sim 1/p$
 - Storage requirements $\sim p$
 - Space/time tradeoff
 - Analytical solution for t^{*1}


Fast Queries

- Batch queries using **MapReduce** and **GFS**
 - F and Q are files in the GFS (with replication)
 - F ~ 64GB, Q ~ 8MB
 - F is stored in GFS chunks
 - Number of mappers = number of F chunks
 - Map:
 - Solves Hamming distance for chunk (64MB) and emits list of near-duplicates
 - Reduce
 - Remove duplicates

Experimental results

- Detecting near-duplicate web pages
 - Web-crawling in Google
- Database
 - 8B fingerprints, $k=1..10$
- Manually tag experimental data set
 - True/false positive/negative
- Precision/recall graph
 - Precision: $\#tp / \# \text{ returned results}$
 - Recall: $\#tp / \# \text{ expected results}$

Experimental results



Papers

- [1] Manku, Gurmeet Singh, Arvind Jain, and Anish Das Sarma. "***Detecting near-duplicates for web crawling.***" In Proceedings of the 16th international conference on World Wide Web, pp. 141-150. ACM, 2007.
- [2] Hoad, Timothy C., and Justin Zobel. "***Methods for identifying versioned and plagiarized documents.***" Journal of the American society for information science and technology 54, no. 3 (2003): 203-215.
- [3] Charikar, Moses S. "***Similarity estimation techniques from rounding algorithms.***" In Proceedings of the thirty-fourth annual ACM symposium on Theory of computing, pp. 380-388. ACM, 2002.
- [4] Henzinger, Monika. "***Finding near-duplicate web pages: a large-scale evaluation of algorithms.***" Proceedings of the 29th annual international ACM SIGIR conference on Research and development in information retrieval. ACM, 2006.